pyconcurrent

Release 2.9.0

Gene C

CONTENTS:

Рy	Python Module Index					
9	Indices and tables					
8		19 19				
	7.3 Our Responsibilities 7.4 Scope	17 18 18 18 18				
7	7.1 Our Pledge	17 17				
6	How to help with this project 6.1 Important resources	15 15 15				
4 5	Changelog 4.1 Tags 4.2 Commits MIT License	9 9 9				
3	Appendix 3.1 Installation	7 7 8 8				
2	Getting Started 2.1 pyconcurrent module	3				
1	pyconcurrent1.1 Overview1.2 Key features1.3 New / Interesting	1 1 1 1				

Index 27

ONE

PYCONCURRENT

1.1 Overview

pyconcurrent is a python class that provides a simple way to do concurrent processing. It supports both asyncio and multiprocessing. The tasks to be run concurrently can either be an executable which is run as a subprocess or a python function to be called.

1.2 Key features

- Provides two classes to do the work: ProcRunAsyncio and ProcRunMp
- Results are provided by the *results* attribute in each class. This is a list of *ProcResults*; one per run.
- Documentation includes the API reference.
- pytest classes validate that all functionality works as it should.

1.3 New / Interesting

- New function run_prog() to run external command. Strictly speaking, this has nothing to do with concurrency, but doing this robustly can be a little tricky. So it is included here.
- PEP 561: Mark module as typed. Now mypy run on code using this module will have the type hints.
- Asyncio now uses the recommended TaskGroup class together with the timeout() context manager. These were introduced in python 3.11. This newer approach is cleaner, more robust and ensures all tasks are appropriately cancelled in the event one task fails. It also offers superior timeout capabilities.
- Timeout now works when using a caller provided function in addition to subprocesses.

GETTING STARTED

All git tags are signed with arch@sapience.com key which is available via WKD or download from https://www.sapience.com/tech. Add the key to your package builder gpg keyring. The key is included in the Arch package and the source= line with *?signed* at the end can be used to verify the git tag. You can also manually verify the signature

2.1 pyconcurrent module

Please see the API reference for additional details.

Here are a couple of simple examples illustrating how the module can be used.

2.1.1 Example 1a: Asyncio

This example uses asyncio and subprocesses to call an executable. tasks must be a list of (key, arg) pairs, 1 per task.

key is a unique identifier, used by caller, one per task. arg is an additional argument for each task; typically arg provides for whatever work that task is responsible for. Each result returned contains both the key and the arg used by that task, information about the success of the task as well as any outputs produced by the task. See ProcResult class for more detail.

This example has 5 tasks to be run concurrently, at most 4 at a time. The results are available in the *proc_run.result*, which is a list of *ProcResult* items; one per task. Since the result order is not pre-defined, each task is identifiable by it's *key* available in the : *result.key*.

```
#!/usr/bin/python
import asyncio
from pyconcurrent import ProcRunAsyncio

async def main():
    """pargs can have additional arguments."""
    pargs = ['/usr/bin/sleep']
    tasks = [(1, 1), (2, 7), (3, 2), (4, 2), (5, 1)]

    proc_run = ProcRunAsyncio(pargs, tasks, num_workers=4, timeout=30)
    await proc_run.run_all()
    proc_run.print_results()

if __name__ == '__main__':
    asyncio.run(main())
```

2.1.2 Example 1b: Multiprocessing

To switch to *multiprocessing* simply replace *ProcRunAsyncio* with *ProcRunMp*, and drop *await* since MP is not *async*. i.e.

```
#!/usr/bin/python
from pyconcurrent import ProcRunMp

def main()
    pargs = ['/usr/bin/sleep']
    tasks = [(1, 1), (2, 7), (3, 2), (4, 2), (5, 1)]

    proc_run = ProcRunMp(pargs, tasks, num_workers=4, timeout=30)
    proc_run.run_all()
    proc_run.print_results()

if __name__ == '__main__':
    main()
```

2.1.3 Example 2: Asnycio

The next example uses a caller supplied function together with asyncio. As in the first example, there are 5 tasks to do and the number of workers is 4, so that 4 tasks are permitted to be run simultaneously.

```
#!/usr/bin/python
import asyncio
from pyconcurrent import ProcRunAsyncio
async def test_func_async(key, args) -> (bool, []):
    """return 2-tuple (success, result)."""
   success = True
                                # pull off the last argument
   nap = args[-1]
    await asyncio.sleep(nap)
    answer = {
            'key' : key,
            'args' : args,
            'success' : success,
            'result' : 'test_func done',
          }
   return (success, answer)
async def main():
   pargs = [test_func_async, 'dummy-arg']
   tasks = [(1, 1), (2, 7), (3, 2), (4, 2), (5, 1)]
    proc_run = ProcRunAsyncio(pargs, tasks, num_workers=4, timeout=30)
    await proc_run.run_all()
   proc_run.print_results()
if __name__ == '__main__':
    asyncio.run(main())
```

For equivalent multiprocessor version for this one, same as above, simply replace ProcRunAsyncio with ProcRunMp

and drop any references to async/await.

The caller supplied function here, $test_func_async()$, must return a 2-tuple of (success:bool, answer:Any) where success should be True if function succeeded.

The function may optionally raise a *RuntimeError* exception, but typically setting *success* is sufficient. If you are using exceptions then please use this one.

2.1.4 Example 3: Non-concurrent

This one shows a non-concurrent external program being executed.

```
#!/usr/bin/python
from pyconcurrent import run_prog
def main()
   pargs_good = ['/usr/bin/sleep', '1']
   pargs_bad = ['/usr/bin/false']
    for pargs in [pargs_good] + [pargs_bad]:
        print(f'Testing: {pargs}:')
        (ret, stdout, stderr) = run_prog(pargs)
        if ret == 0:
            print('\tAll well')
            print(stdout)
        else:
            print('\tFailed')
            print(stderr)
if __name__ == '__main__':
   main()
```

THREE

APPENDIX

3.1 Installation

Available on

- Github
- · Archlinux AUR

On Arch you can build using the provided PKGBUILD in the packaging directory or from the AUR. All git tags are signed with arch@sapience.com key which is available via WKD or download from https://www.sapience.com/tech. Add the key to your package builder gpg keyring. The key is included in the Arch package and the source= line with *?signed* at the end can be used to verify the git tag. You can also manually verify the signature

```
git tag -v <tag-name>
```

To build manually, clone the repo and:

```
rm -f dist/*
/usr/bin/python -m build --wheel --no-isolation
root_dest="/"
./scripts/do-install $root_dest
```

When running as non-root then root_dest must be a user writable directory

3.2 Dependencies

Run Time:

• python (3.13 or later)

Building Package:

- git
- hatch (aka python-hatch)
- wheel (aka python-wheel)
- build (aka python-build)
- installer (aka python-installer)
- rsync
- pytest (aka python-pytest)
- pytest-asyncio (aka python-pytest-asyncio)

Optional for building docs:

- sphinx
- myst-parser (aka python-myst-parser)
- sphinx-autoapi (aka python-sphinx-autoapi)
- texlive-latexextra (archlinux packaging of texlive tools)

3.3 Philosophy

We follow the *live at head commit* philosophy. This means we recommend using the latest commit on git master branch. We also provide git tags.

This approach is also taken by Google¹².

3.4 License

Created by Gene C. and licensed under the terms of the MIT license.

- SPDX-License-Identifier: MIT
- SPDX-FileCopyrightText: © 2025-present Gene C <arch@sapience.com>

¹ https://github.com/google/googletest

² https://abseil.io/about/philosophy#upgrade-support

FOUR

CHANGELOG

4.1 Tags

```
1.1.2 (2025-04-24) -> 2.9.0 (2025-06-26)
48 commits.
```

4.2 Commits

• 2025-06-26 : **2.9.0**

```
run_prog: Make sure iolists are updated before calling select()
2025-06-22 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf
```

• 2025-06-22 : **2.8.0**

run_prog - make internal helper class internal so not included in API ref
update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-06-22 : **2.7.0**

```
run_prog: refactor code. Protect against one possible deadlock
2025-06-19 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf
```

• 2025-06-19 : **2.6.0**

run_prog: When writing to subprocess stdin, add a flush() before we close the file object update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-06-19 : **2.5.0**

```
run_prog: make stdin nonblock as well
2025-06-17 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf
```

• 2025-06-17 : 2.4.1

```
PEP-484 use builtin tuple instead of Tuple update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf
```

• 2025-06-17 : **2.4.0**

New function run_prog() to run external command.
2025-05-21 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-05-21 : **2.3.0**

Use builtin types where possible. e.g. typing.List -> list 2025-05-19 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-05-19 : **2.2.3**

Arch PKGBUILD: move pytest dependency **from makedepends** to checkdepends. 2025-05-03 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-05-03 : **2.2.2**

Remove old test directory update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-05-03 : **2.2.1**

pytests: reorganize a little. tests are unchanged
2025-05-01 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-05-01 : **2.2.0**

PEP 561 type hints (module users using mypy etc get type hints)
Update status to Production/Stable
2025-04-30 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-30 : **2.1.0**

Bug fix returning stdout string **from subprocess**2025-04-27 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-27 : **2.0.2**

API ref rst format
Fix typo in api reference

• 2025-04-27 : **2.0.0**

Asyncio now uses the recommended TaskGroup class together with
the timeout() context manager. These were introduced in python 3.11.
This newer approach is cleaner, more robust and ensures all tasks
are appropriately cancelled in the event one task fails. It also offers
superior timeout capabilities.
Timeout now works when using a caller provided function in addition to
subprocesses.

2025-04-26 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-26 : **1.5.0**

Style per PEP-8, PEP-257 and PEP-484.
Bug fix with stdout/stderr returned for asyncio using subprocesses update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-26 : **1.4.1**

More consistent with PEP-8 & PEP-257.

Primarily double quotes for docstrings instead of single and ensure ends with period

2025-04-25 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-25 : **1.4.0**

asyncio : missing decode() for stdout/err

pytest : ensure python path uses dev source for tests

Add missing SPDX identifiers

update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-25 : **1.3.3**

Fix readme typo **and** small tweak **for** clarity

2025-04-24 update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-24 : **1.3.2**

Change examples in README to include everything to actually run update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-24 : **1.3.1**

Add note about git signing key in readme update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-24 : **1.3.0**

Add missing tests dir after it was moved update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-24 : **1.2.0**

Move tests dir to top level update Docs/Changelogs Docs/_build/html Docs/pyconcurrent.pdf

• 2025-04-24 : **1.1.3**

Add dateutil dep to PKGBUILD

• 2025-04-24 : **1.1.2**

Initial Commit

4.2. Commits

FIVE

MIT LICENSE

Copyright © 2025-present Gene C <arch@sapience.com>

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

HOW TO HELP WITH THIS PROJECT

Thank you for your interest in improving this project. This project is open-source under the MIT license.

6.1 Important resources

• Git Repo

6.2 Reporting Bugs or feature requests

Please report bugs on the issue tracker in the git repo. To make the report as useful as possible, please include

- · operating system used
- · version of python
- explanation of the problem or enhancement request.

6.3 Code Changes

If you make code changes, please update the documentation if it's appropriate.

CONTRIBUTOR COVENANT CODE OF CONDUCT

7.1 Our Pledge

In the interest of fostering an open and welcoming environment, we as contributors and maintainers pledge to making participation in our project and our community a harassment-free experience for everyone, regardless of age, body size, disability, ethnicity, sex characteristics, gender identity and expression, level of experience, education, socio-economic status, nationality, personal appearance, race, religion, or sexual identity and orientation.

7.2 Our Standards

Examples of behavior that contributes to creating a positive environment include:

- Using welcoming and inclusive language
- · Being respectful of differing viewpoints and experiences
- · Gracefully accepting constructive criticism
- · Focusing on what is best for the community
- Showing empathy towards other community members

Examples of unacceptable behavior by participants include:

- The use of sexualized language or imagery and unwelcome sexual attention or advances
- Trolling, insulting/derogatory comments, and personal or political attacks
- · Public or private harassment
- Publishing others' private information, such as a physical or electronic address, without explicit permission
- Other conduct which could reasonably be considered inappropriate in a professional setting

7.3 Our Responsibilities

Maintainers are responsible for clarifying the standards of acceptable behavior and are expected to take appropriate and fair corrective action in response to any instances of unacceptable behavior.

Maintainers have the right and responsibility to remove, edit, or reject comments, commits, code, wiki edits, issues, and other contributions that are not aligned to this Code of Conduct, or to ban temporarily or permanently any contributor for other behaviors that they deem inappropriate, threatening, offensive, or harmful.

7.4 Scope

This Code of Conduct applies both within project spaces and in public spaces when an individual is representing the project or its community. Examples of representing a project or community include using an official project e-mail address, posting via an official social media account, or acting as an appointed representative at an online or offline event. Representation of a project may be further defined and clarified by project maintainers.

7.5 Enforcement

Instances of abusive, harassing, or otherwise unacceptable behavior may be reported by contacting the project team at <arch@sapience.com>. All complaints will be reviewed and investigated and will result in a response that is deemed necessary and appropriate to the circumstances. The Code of Conduct Committee is obligated to maintain confidentiality with regard to the reporter of an incident. Further details of specific enforcement policies may be posted separately.

7.6 Attribution

This Code of Conduct is adapted from the Contributor Covenant, version 1.4, available at https://www.contributor-covenant.org/version/1/4/code-of-conduct.html

7.7 Interpretation

The interpretation of this document is at the discretion of the project team.

EIGHT

API REFERENCE

This page contains auto-generated API reference documentation¹.

8.1 pyconcurrent

Public Methods. pyconcurrent.

8.1.1 Submodules

pyconcurrent.proc_asyncio

Concurrent tasks using asyncio.

Module Contents

class ProcRunAsyncio(pargs: list[Any], tasks_todo: list[tuple[Any, Any]], num_workers: int = 4, timeout: int = 0, verb: bool = False)

Run concurrent processes using asyncio.

Asynio concurrent process runs. Supports program to be run as a subprocess or a function to be called. The result of each run is returned as in ProcResult class instance.

Args:

pargs ([Any]):

The first element is the command/function to be run and remainder are any additional arguments.

tasks_todo ([(Any, Any)]):

list of task items to be run concurrently. Each task is a 2-tuple, (key, arg).

- key is a unique identifier, converable to string via str(key)
- arg is an additional argument to the routine when it is called.

Both key and arg are saved into the result class instance returned.

num workers (int):

Max number of processes to use. Value of 0 is unlimited and 1 will mean each is run serially one at a time.

timeout (int):

The maximum number of seconds allotted to each process. If not complete within "timeout", then process/function will be cancelled/killed and the "result" instance will include:

• res.success set to False

¹ Created with sphinx-autoapi

```
• res.timeout set to True.
           verb (bool):
               If set to true, some additional information is sent to stdout.
      Attributes:
           result ([ProcResult]):
               list of results, one per task. See ProcResult for more detail.
      Methods:
      print_results()
           Test tool: prints each result using the ProcResul::print().
      async run_all()
           Start running all the provided commands/functions concurrently.
           Awaitable, so caller is responsible for calling asyncio.run(). See run_all_start_asyncio() for non-awaitable
           version.
pyconcurrent.proc_mp
Concurrent tasks using multiprocessing.
Module Contents
 \textbf{class ProcRunMp}(pargs: list[Any], tasks\_todo: list[tuple[Any, Any]], num\_workers: int = 4, timeout: int = 0, \\
                    verb: bool = False)
      Run concurrent processes using multiprocessing.
      Same calling convention as ProcRunAsyncio.
      Note: func cannot be async func() - conflicts with mp starmap using async
      print_results()
           Test tool: prints each result using the ProcResul::print().
      run_all()
           Do the work
pyconcurrent.proc result
Result class which ProcRunMp/ProcRunAsyncio use.
Module Contents
class ProcResult(key, arg)
      Result of running one of the concurrent processes.
      Args:
           key (Any):
               Caller provided unique identifier.
           arg (Any):
```

The additional argument used for this run.

Attributes:

```
time start (float):
                                         Unix time in seconds.
                             time run (float):
                                         Seconds taken for this item to complete.
                             success (bool):
                                         True if completed successfully.
                             timeout (bool):
                                         True if failed to complete in less than timeout restriction.
                             key (Any):
                                         The caller provided unique identifier.
                             arg (Any):
                                         The called provided argument for this run.
                             returncode (int):
                                         Return value of subprocess. Typically 0 for success.
                                         Returned stdout of subprocess.
                             stderr (str):
                                         Returned stderr of subprocess.
                             answer (Any):
                                         Return provided by the function.
               print()
                             Testing: simple print attributes.
pyconcurrent.run_prog
External program execution
Module Contents
run\_prog(pargs: list[str], input\_str: str | None = None, stdout: int = subprocess.PIPE, stderr: int =
                             subprocess.PIPE, env: dict[str, str] \mid None = None, test: bool = False, verb: bool = False) \rightarrow tuple[int, str] \mid None = None, test: bool = False, verb: bool = False) \rightarrow tuple[int, str] \mid None = None, test: bool = False, verb: bool = False) \rightarrow tuple[int, str] \mid None = None, test: bool = False, verb: bool = False) \rightarrow tuple[int, str] \mid None = None, test: bool = False, verb: bool = False) \rightarrow tuple[int, str] \mid None = None, test: bool = False, verb: bool = False) \rightarrow tuple[int, str] \mid None = None, test: bool = False, verb: bool = False) \rightarrow tuple[int, str] \mid None = None, test: bool = False, verb: bool = False, v
                             str, str]
               Run external program using subprocess.
               Take care to handle large outputs (default buffer size is 8k). This avoids possible hangs should IO buffer fill up.
               non-blocking IO together with select() loop provides a robust methodology.
               Args:
                             pargs (list[str]):
                                         The command + arguments to be run in standard list format. e.g. ['/usr/bin/sleep', '22'].
                             input str (str | None):
                                         Optional input to be fed to subprocess stdin. Defaults to None.
                             stdout (int):
                                         Subprocess stdout. Defaults to subprocess.PIPE
```

8.1. pyconcurrent 21

Subprocess stderr. Defaults to subprocess.PIPE

env (None | dict[str, str]):

Optional to specify environment for subprocess to use. If not set, inherits from calling process as usual.

test (bool):

Flag - if true dont actually run anything.

verb (bool):

Flag - only used with test == True - prints pargs.

Returns:

tuple[retc: int, stdout: str, stderr: str]:

retc is 0 when all is well. stdout is what the subprocess returns on it's stdout and stderr is what it's stderr return.

Note that any input string is written in it's entirety in one shot to the subprocess. This should not be a problem.

pyconcurrent.version

Project pyconcurrent.

Module Contents

 $version() \rightarrow str$

report version and release date

NINE

INDICES AND TABLES

- genindex
- modindex
- search

PYTHON MODULE INDEX

p

```
pyconcurrent, 19
pyconcurrent.proc_asyncio, 19
pyconcurrent.proc_mp, 20
pyconcurrent.proc_result, 20
pyconcurrent.run_prog, 21
pyconcurrent.version, 22
```

26 Python Module Index

INDEX

```
M
module
    pyconcurrent, 19
    pyconcurrent.proc_asyncio, 19
    pyconcurrent.proc_mp, 20
    {\tt pyconcurrent.proc\_result}, 20
    pyconcurrent.run_prog, 21
    pyconcurrent.version, 22
Р
print() (ProcResult method), 21
print_results() (ProcRunAsyncio method), 20
print_results() (ProcRunMp method), 20
ProcResult (class in pyconcurrent.proc_result), 20
ProcRunAsyncio (class in pyconcurrent.proc_asyncio),
ProcRunMp (class in pyconcurrent.proc_mp), 20
pyconcurrent
    module, 19
pyconcurrent.proc_asyncio
    module, 19
pyconcurrent.proc_mp
    module, 20
pyconcurrent.proc_result
    module, 20
pyconcurrent.run_prog
    module, 21
pyconcurrent.version
    module, 22
R
run_all() (ProcRunAsyncio method), 20
run_all() (ProcRunMp method), 20
run_prog() (in module pyconcurrent.run_prog), 21
V
version() (in module pyconcurrent.version), 22
```